

Spring Boot + Core Java Interview Bootcamp

5-Day Crash Course for the Java Microservices Developer Interview
Complete Master Edition — Modules 1–11

Targeting interviews at Google · Amazon · Microsoft · Uber · Stripe · Netflix · Oracle · VMware · TCS · Infosys · Cognizant ·
Accenture · Capgemini · IBM · Deloitte · EY · Wipro · LTIMindtree

Every topic: why it's asked · internals · ASCII memory diagrams · production examples · traps · best answers ·
code · follow-ups · cheat sheets

Generated July 03, 2026

Contents

Module 1 — Core Java

Module 2 — Collections Framework

Module 3 — Java 8+ (Lambdas, Functional Interfaces, Streams, Optional)

Module 1 — Core Java

Highest priority. Rounds 1 and 2 at every service company (TCS, Infosys, Cognizant, Accenture, Capgemini, Wipro, LTIMindtree) and the “Java fundamentals” screen at product companies live here. If you are shaky on JVM memory, OOP, and `String`, you will not pass.

Node.js bridge: You already know V8, the event loop, and `require`. Java’s JVM is V8’s much older, heavier cousin: explicit typing, ahead-of-time byte-code compilation, real OS threads (not a single event loop), and a garbage collector you can actually tune.

1.1 JVM, JDK, JRE & the Compilation Process

1. Why Interviewers Ask This

It’s the fastest way to tell a “Java developer” from someone who copied Spring code. If you can’t cleanly separate JDK/JRE/JVM you’ll be filtered in the first 5 minutes.

2. Core Concept

- **JDK** (Java Development Kit) = JRE + development tools (`javac`, `jar`, `javadoc`, `jdb`). You need it to *build*.
- **JRE** (Java Runtime Environment) = JVM + core libraries (`java.lang`, `java.util`...). You need it to *run*.
- **JVM** (Java Virtual Machine) = the abstract machine that executes byte-code. It’s a *specification*; HotSpot is Oracle’s implementation.

```
JDK  ⊃  JRE  ⊃  JVM
(build) (run) (execute bytecode)
```

3. Internal Working — Compilation & Execution

```
Person.java --javac--> Person.class (bytecode) --> JVM ClassLoader
                                     |
                                     +-----+-----+
                                     | Bytecode Verifier |
                                     +-----+-----+
                                     |
                               Interpreter (fast start) + JIT compiler
                                     |
                               native machine code (cached)
```

- `javac` compiles source to **platform-independent byte-code** (`.class`).
- JVM **interprets** byte-code initially (fast startup).
- The **JIT (Just-In-Time) compiler** watches for **hot** methods (default threshold ~10,000 invocations) and compiles them to native code. HotSpot has two JITs: **C1** (client, quick) and **C2** (server, aggressive optimizations); modern JVMs use **tiered compilation** (C1 first, then C2).
- “Write once, run anywhere” = byte-code is portable; only the JVM is platform-specific.

4. Memory Diagram

```
+-----+-----+ JVM Process +-----+-----+
| ClassLoader Subsystem -> Runtime Data Areas -> Execution Engine |
|                               (Heap, Stacks,                        (Interpreter,
|                               Metaspace, PC,                        JIT, GC)
|                               Native Method Stack)                  |
+-----+-----+-----+-----+
```

5. Real Production Example

On a Spring Boot service you ship a fat JAR. Build server needs the **JDK**; the container base image (`eclipse-temurin:21-jre`) only needs the **JRE** — smaller, fewer CVEs. JIT “warm-up” is why the *first* few requests after deploy are slow and p99 latency spikes right after a rolling restart.

6. Most Asked Interview Questions

- Difference between JDK, JRE, JVM? (*follow-up: which do you put in a Docker prod image?*)
- Is Java compiled or interpreted? (*both — bytecode compiled, then interpreted + JIT-compiled*)
- What is JIT? What is tiered compilation? (*follow-up: what is warm-up?*)
- Why is Java platform independent but the JVM is not?

7. Interview Traps

- Saying “Java is purely interpreted” (wrong — JIT compiles to native).
- Saying byte-code is machine code (it's not — it's JVM instructions).
- Thinking the JRE can compile code (it can't; no `javac`).

8. Best Answer

“`javac` compiles `.java` into portable byte-code. At runtime the JVM's class loader loads and verifies it, the interpreter runs it immediately for fast startup, and the JIT compiler promotes hot methods to native code using tiered C1/C2 compilation. The JDK is what I build with; production containers only ship a JRE to stay small and secure.”

9. Coding Example

```
// javac Hello.java -> Hello.class ; java Hello
public class Hello {
    public static void main(String[] args) {
        System.out.println("Bytecode ran on " + System.getProperty("java.version"));
    }
}
// Inspect bytecode: javap -c Hello
```

10. Follow-up Coding Questions

- Run `javap -c` on a class and explain a few opcodes (`iload`, `invokevirtual`).
- How would you force C2-only or disable tiered compilation? (`-XX:-TieredCompilation`)

11. Summary

JDK builds, JRE runs, JVM executes. Byte-code is portable; JIT makes it fast.

12. Cheat Sheet

Term	Contains	Purpose
JDK	JRE + javac/jar/javadoc	Development/build
JRE	JVM + libraries	Run apps
JVM	Class loader + memory + engine	Execute bytecode
JIT	C1 + C2 (tiered)	Bytecode → native

1.2 Class Loading

1. Why Interviewers Ask This

`ClassNotFoundException` vs `NoClassDefFoundError`, and how Spring/Tomcat isolate apps, all come from class loading. Frequent at product companies and any “we had a jar hell bug” story.

2. Core Concept

Classes are loaded lazily, on first active use, by a hierarchy of class loaders using **delegation**.

3. Internal Working — Loading → Linking → Initialization

```
1. Loading      : find .class bytes, create Class object in Metaspace
2. Linking
  a. Verification : bytecode is valid & safe
  b. Preparation  : static fields get default values (0/null/false)
  c. Resolution   : symbolic refs -> direct refs
3. Initialization : run static blocks & static initializers (assign real values)
```

Parent-delegation model: a loader asks its *parent* first; only if the parent can't find the class does the child try. Prevents core classes from being spoofed.

```
Bootstrap (rt/core, native) -> loads java.* (parent of all)
  ^
Platform/Extension loader  -> loads platform modules
  ^
Application/System loader  -> loads your classpath classes
  ^
(Custom loaders: web apps, plugins, OSGi)
```

4. Memory Diagram

```
new Foo() -> AppClassLoader.loadClass("Foo")
           -> ask Platform loader -> ask Bootstrap loader
           Bootstrap: "not mine" Platform: "not mine"
           App loader loads Foo.class -> Metaspace [Class<Foo>]
```

5. Real Production Example

Two web apps in one Tomcat each get their own `WebAppClassLoader`, so both can use different Jackson versions without clashing. `NoClassDefFoundError` in prod usually = the class was present at compile time but a dependency/jar is missing at runtime (classpath mismatch).

6. Most Asked Interview Questions

- Explain the class loading process. (3 phases)
- What is parent delegation and why? (security + no duplicates)
- `ClassNotFoundException` vs `NoClassDefFoundError`? (former: `Class.forName` can't find at runtime; latter: was there at compile, gone/failed init at runtime)
- When does a static block run? (at initialization, once, thread-safely)

7. Interview Traps

- Confusing the two class errors (very common).
- Thinking loading a class also creates instances (it doesn't).
- Saying static fields get real values in *preparation* (no — defaults there, real values in *initialization*).

8. Best Answer

“Loading finds the bytes and builds the `Class` object; linking verifies, gives static fields defaults, and resolves references; initialization runs static initializers once. Loaders delegate to their parent first so core `java.*` classes can't be overridden. That isolation is how one Tomcat runs two apps with conflicting library versions.”

9. Coding Example

```
public class Loaders {
    static { System.out.println("static init runs once, at initialization"); }
    public static void main(String[] a) {
        System.out.println(String.class.getClassLoader()); // null = Bootstrap
        System.out.println(Loaders.class.getClassLoader()); // AppClassLoader
    }
}
```

10. Follow-up Coding Questions

- Write a custom `ClassLoader` that loads a class from a byte array.
- Demonstrate lazy loading: prove a class isn't initialized until first active use.

11. Summary

Load → Link (verify, prepare, resolve) → Initialize. Parents get first refusal.

12. Cheat Sheet

Error	When	Cause
<code>ClassNotFoundException</code>	reflection/ <code>Class.forName</code>	name not on classpath
<code>NoClassDefFoundError</code>	linking/init	present at compile, missing/failed at runtime

1.3 Java Memory Model — Heap, Stack, Metaspace

1. Why Interviewers Ask This

Memory questions separate juniors from seniors and lead straight into GC and `OutOfMemoryError` debugging — a top production topic.


2. Core Concept

- **Heap** — shared, GC-managed; all objects & arrays live here.
- **Stack** — per-thread; holds frames with local variables and references (LIFO).
- **Metaspace** — off-heap (native memory since Java 8); class metadata, replaced PermGen.
- **PC register & Native method stack** — per thread, small.

3. Internal Working

- Each thread gets its own **stack**; each method call pushes a **frame** (locals, operand stack, return address). Frame pops on return. Deep recursion → `StackOverflowError`.
- **Heap** is split for generational GC: **Young** (Eden + two Survivor spaces) and **Old/Tenured**.
- **Metaspace** grows in native memory (bounded by `-XX:MaxMetaspaceOnly` if set); before Java 8 this was **PermGen** inside the heap and caused frequent `OutOfMemoryError: PermGen space`.
- **Primitives declared as locals** live on the stack; **object fields** (even primitive fields) live on the heap inside the object.

4. Memory Diagram



5. Real Production Example

A memory leak in a Spring app = objects unintentionally reachable (e.g., a static `Map` cache never evicted) → Old gen fills → `OutOfMemoryError: Java heap space`. You capture a heap dump (`-XX:+HeapDumpOnOutOfMemoryError`) and analyze in Eclipse MAT. `-Xmx2g -Xms2g` fixes heap size; classloader leaks (redeploys) show as Metaspace growth.

6. Most Asked Interview Questions

- Stack vs heap? Where do primitives/objects/references live? (follow-up: are *String literals* on stack? No — heap string pool)
- What replaced PermGen and why? (Metaspace, native memory, auto-grows)
- What causes `StackOverflowError` vs `OutOfMemoryError`?
- Is memory allocation thread-safe on the heap? (yes; TLABs make it fast)

7. Interview Traps

- “Objects can live on the stack” — no (ignoring escape-analysis scalar replacement, which is an optimization detail).
- “Primitives always on stack” — only *local* primitives; primitive *fields* are on the heap in the object.
- Confusing Metaspace (native) with heap.

8. Best Answer

“Objects live on the shared heap, split into young and old generations for GC. Each thread has its own stack of frames holding locals and references — references point into the heap. Class metadata moved from PermGen to native Metaspace in Java 8, which auto-grows so we stopped seeing PermGen OOMs. Heap OOM means a leak or undersized `-Xmx`; `StackOverflow` means unbounded recursion.”

9. Coding Example

```

public class MemoryDemo {
    int fieldOnHeap = 10;           // lives in the object on the heap
    void method() {
        int localOnStack = 5;      // stack
        MemoryDemo ref = new MemoryDemo(); // 'ref' on stack, object on heap
    }
    static int recurse(int n){ return recurse(n+1); } // -> StackOverflowError
}

```

10. Follow-up Coding Questions

- Trigger and catch a `StackOverflowError`; why can you catch an `Error` but shouldn't rely on it?
- Write code that causes a heap OOM with an ever-growing `List`.

11. Summary

Heap = objects (shared, GC). Stack = per-thread frames. Metaspace = class metadata (native).

12. Cheat Sheet

Area	Scope	Stores	Error
Heap	shared	objects/arrays	<code>OutOfMemoryError: Java heap space</code>
Stack	per-thread	frames, locals, refs	<code>StackOverflowError</code>
Metaspace	shared (native)	class metadata	<code>OutOfMemoryError: Metaspace</code>

1.4 Garbage Collection, Generational GC & Object Lifecycle

1. Why Interviewers Ask This

GC is *the* differentiator for backend Java. Everyone from TCS to Netflix asks “how does GC work” and “which collector do you use and why”.

2. Core Concept

GC automatically reclaims unreachable objects. **Reachability** (from GC roots), not reference counting, decides liveness. Generational GC exploits the **weak generational hypothesis**: *most objects die young*.

3. Internal Working — Generational GC

```
1. New objects -> Eden.
2. Eden full -> Minor GC: live objects copied to a Survivor space; dead ones dropped.
3. Survivors aged each Minor GC. After threshold (MaxTenuringThreshold) -> promoted to Old gen.
4. Old gen full -> Major/Full GC (slower, may 'stop the world').
```

- **GC Roots**: stack locals, static fields, active threads, JNI refs. Anything reachable from a root is live.
- **Mark-Sweep-Compact**: mark live, sweep dead, compact to remove fragmentation.
- **Collectors**:
- **Serial** — single thread, tiny heaps.
- **Parallel (Throughput)** — multi-thread, batch jobs.
- **G1 (default since Java 9)** — region-based, predictable pause targets (`-XX:MaxGCPauseMillis`).
- **ZGC / Shenandoah** — sub-millisecond pauses, huge heaps (low-latency services).
- **Stop-The-World (STW)**: application threads pause during certain GC phases. Minimizing STW is the whole game.

4. Memory Diagram

```
Allocation: [ Eden ] fills -> Minor GC -> survivors -> [ S0/S1 ] -> age -> [ Old ]
GC roots ---> reachable objects kept; unreachable islands (even if they reference
each other) are collected. Cyclic refs are NOT a leak in Java.
```

5. Real Production Example

A latency-sensitive payment service on a 8GB heap uses **G1** with `-XX:MaxGCPauseMillis=200`. You watch GC logs (`-Xlog:gc*`) for frequent Full GCs (a red flag for leaks or undersized old gen). Netflix-style low-latency services move to **ZGC** to keep pauses under 1ms.

6. Most Asked Interview Questions

- How does GC decide what to collect? (*reachability from GC roots, not ref counting*)
- Explain generational GC / why young & old? (*most objects die young → cheap minor GC*)
- Can you force GC? (`System.gc()` is only a hint; never rely on it)
- Difference between Minor, Major, Full GC? What is Stop-The-World?
- Which collectors do you know? When G1 vs ZGC? (*follow-up: does GC prevent all memory leaks? No.*)
- Do circular references cause leaks in Java? (*No — reachability handles cycles.*)

7. Interview Traps

- Saying Java uses reference counting (it doesn't — reachability).
- Claiming `System.gc()` guarantees collection.
- Saying "Java can't leak memory" — it can (unbounded caches, listeners, ThreadLocals, classloaders).
- Confusing `finalize()` with a destructor (see next section).

8. Best Answer

"The GC frees objects unreachable from GC roots — so cycles aren't a problem. It's generational because most objects die young: cheap minor GCs sweep Eden, survivors age and get promoted to old gen, which is collected less often by an expensive major GC. Modern JVMs default to G1, which is region-based and lets me target a max pause time; for ultra-low latency I'd use ZGC. `System.gc()` is only a hint, and GC doesn't fix logical leaks like an unbounded static cache."

9. Coding Example

```
public class GcDemo {
    public static void main(String[] args) {
        Object a = new Object();
        a = null; // old object now unreachable -> eligible for GC
        System.gc(); // hint only; do NOT rely on it in code
        // Prefer WeakReference/soft caches to let GC reclaim under pressure:
        java.lang.ref.WeakReference<byte[]> cache =
            new java.lang.ref.WeakReference<>(new byte[1024]);
        System.out.println(cache.get() != null);
    }
}
```

10. Follow-up Coding Questions

- Write a class that leaks via a `static List` and explain how to detect it (heap dump/MAT).
- Explain `WeakReference` vs `SoftReference` vs `PhantomReference` with a cache example.

11. Summary

Reachability-based, generational, mostly automatic. G1 default; ZGC for low pause. GC ≠ leak-proof.

12. Cheat Sheet

Collector	Best for	Pause
Serial	tiny/CLI	high
Parallel	throughput/batch	medium
G1 (default)	general services	tunable (~200ms)
ZGC/Shenandoah	low latency, big heap	<1ms

Object lifecycle: Created (`new`) → In use (reachable) → Unreachable → Eligible → (optional `finalize`) → Collected.

1.5 OOP, SOLID & the Four Pillars

1. Why Interviewers Ask This

Every design and code-quality question rests on OOP + SOLID. "Explain SOLID with an example" is nearly guaranteed in service-company and mid-level product interviews.

2. Core Concept — Four Pillars

- **Encapsulation** — bundle state + behavior; hide internals behind methods (private fields, getters/setters, invariants).
- **Abstraction** — expose *what*, hide *how* (interfaces, abstract classes).
- **Inheritance** — reuse via “is-a” (`Dog extends Animal`).
- **Polymorphism** — one interface, many forms: **compile-time** (overloading) & **runtime** (overriding via dynamic dispatch).

3. Internal Working — Runtime Polymorphism

The JVM keeps a **vtable** (virtual method table) per class. An overridden call (`animal.sound()`) is resolved at runtime by the *actual object's* type via `invokevirtual`, not the reference type. Overloading is resolved at *compile time* by argument types via `invokestatic/invokevirtual` with a fixed signature.

4. Memory Diagram

```
Animal a = new Dog();           // ref type Animal, object type Dog
a.sound();
    -> invokevirtual -> look up Dog's vtable -> Dog.sound()    (runtime dispatch)
```

5. SOLID (with production framing)

Letter	Principle	One-liner	Spring example
S	Single Responsibility	one reason to change	Controller ≠ Service ≠ Repository
O	Open/Closed	open to extend, closed to modify	add a new <code>PaymentStrategy</code> impl, don't edit existing
L	Liskov Substitution	subtypes usable as base type	<code>Square extends Rectangle</code> violates it
I	Interface Segregation	many small interfaces > one fat	split <code>Repository</code> not one god-interface
D	Dependency Inversion	depend on abstractions	inject <code>PaymentGateway</code> interface, not <code>StripeClient</code>

6. Most Asked Interview Questions

- Explain the 4 OOP pillars with real examples.
- Explain SOLID; give a violation and the fix. (*follow-up: which SOLID does DI implement? → D*)
- Compile-time vs runtime polymorphism? (*overloading vs overriding*)
- Is Java 100% OOP? (*No — primitives exist.*)

7. Interview Traps

- Confusing abstraction with encapsulation (abstraction = design/hide complexity; encapsulation = data hiding + bundling).
- Saying overloading is runtime polymorphism (it's compile-time).
- Reciting SOLID with no example — interviewers want the *fix*.

8. Best Answer

“Encapsulation hides state behind methods to protect invariants; abstraction exposes intent through interfaces; inheritance reuses via is-a; polymorphism lets one reference behave as many types — overloading resolved at compile time, overriding at runtime via the vtable. SOLID keeps that flexible: e.g., Dependency Inversion is exactly what Spring DI gives us — my service depends on a `PaymentGateway` interface, and Spring injects Stripe or a mock.”

9. Coding Example

```
interface PaymentGateway { boolean charge(long cents); }           // abstraction + D + 0

final class StripeGateway implements PaymentGateway {              // encapsulated impl
    public boolean charge(long cents) { /* call Stripe */ return true; }
}

class CheckoutService {                                           // S: only checkout logic
    private final PaymentGateway gateway;                          // D: depends on abstraction
    CheckoutService(PaymentGateway gateway) { this.gateway = gateway; }
    boolean pay(long cents) { return gateway.charge(cents); }
}
```

10. Follow-up Coding Questions

- Refactor an `if/else` on payment type into the Strategy pattern (Open/Closed).
- Show a Liskov violation with `Rectangle/Square` and fix it.

11. Summary

4 pillars + SOLID = maintainable code. DI = Dependency Inversion in practice.

12. Cheat Sheet

Single · **O**pen-closed · **L**iskov · **I**nterface-seg · **D**ependency-inversion. Overload = compile-time; Override = runtime (vtable).

1.6 Interface vs Abstract Class · Overloading vs Overriding

Interface vs Abstract Class

	Interface	Abstract Class
Multiple inheritance	Yes (implement many)	No (extend one)
State/fields	only <code>public static final</code> constants	instance fields allowed
Methods	<code>abstract</code> + <code>default/static/private</code> (Java 8+)	<code>abstract</code> + <code>concrete</code>
Constructor	No	Yes
Use when	capability/contract ("can-do": <code>Comparable</code>)	shared base + state ("is-a")

Trap: Since Java 8, interfaces have `default` methods → the "interfaces can't have bodies" answer is outdated. The **diamond problem** with two default methods is resolved by forcing the class to override and call `Interface.super.method()`.

Overloading vs Overriding

	Overloading	Overriding
When	compile-time	runtime
Signature	must differ (params)	must match
Return type	can differ	same/covariant
Access	any	can't reduce visibility
<code>static</code>	can overload	can't override (hidden, not overridden)

Best answer for "can you override a static method?" → No. A static method belongs to the class; redeclaring it in a subclass **hides** it (resolved by reference type, not object type).

1.7 final · finally · finalize · static · this · super

final, finally, finalize (a guaranteed question)

- **final** — keyword. **final** variable = constant; **final** method = can't override; **final** class = can't extend (e.g., `String`).
- **finally** — block that always runs after **try/catch** (even on return/exception), used for cleanup. Skipped only on `System.exit()` or JVM crash.
- **finalize()** — Object method the GC *might* call before reclaiming. **Deprecated since Java 9, removed in later releases.** Unpredictable, can resurrect objects, hurts GC, no guarantee it runs. Use **try-with-resources** / `Cleaner` / `AutoCloseable` instead.

Best answer: “**final** is a modifier, **finally** is a cleanup block, **finalize()** is a deprecated GC hook I never use — I use try-with-resources for deterministic cleanup.”

static

Belongs to the **class**, not instances: shared across all objects, loaded at class initialization. Static methods can't use **this/super** or access instance members directly. Static blocks run once at init. Watch for the trap: static fields + concurrency = shared mutable state (needs synchronization).

this VS super

- **this** — current object (disambiguate fields, chain constructors `this(...)`).
- **super** — parent (call parent constructor `super(...)`, or parent method `super.foo()`).
- `super(...)/this(...)` must be the **first** statement in a constructor.

1.8 String Pool · String vs StringBuilder vs StringBuffer · Immutability

1. Why Interviewers Ask This

`String` is the most-used class and the classic immutability/`==` vs `.equals()` trap. Nearly guaranteed.

2 & 3. Core + Internal

- `String` is **immutable** and **final**. Its `char[]/byte[]` (compact strings since Java 9) is private and never mutated.
- **String Pool** (a.k.a. intern pool) lives in the **heap** (since Java 7; was PermGen before). String *literals* are interned — identical literals share one object. `new String("x")` creates a *new* heap object (not pooled unless you call `.intern()`).

4. Memory Diagram

```
String a = "hi";           // pool: ["hi"] <- a
String b = "hi";           // b -> same pooled "hi"
String c = new String("hi"); // heap: [new "hi"] <- c
                           //
String d = c.intern();      // d -> pooled "hi"
```

a == b	-> true
a == c	-> false
a.equals(c)	-> true
a == d	-> true

5. Real Production Example

Building a big log/CSV line by `s += token` in a loop creates a new `String` each iteration → $O(n^2)$ garbage. Use **StringBuilder** (single mutable buffer). In a multithreaded logger you'd use **StringBuffer** (synchronized) — or better, keep the builder thread-local.

Immutability — why `String` is immutable

Security (used in class loading, network, file paths), thread-safety (shareable without locks), safe as **HashMap keys** (cached `hashCode`), and pool sharing. Any “modification” returns a **new** `String`.

String vs StringBuilder vs StringBuffer

	String	StringBuilder	StringBuffer
Mutable	No	Yes	Yes
Thread-safe	Yes (immutable)	No	Yes (synchronized)
Speed	slow for concat loops	fastest	slower (locking)
Use	fixed text, keys	single-thread build	rare, shared build

6. Most Asked Interview Questions

- Why is `String` immutable? (security, thread-safety, hashcode caching, pool)
- `==` vs `.equals()` for strings? (reference vs value)
- `new String("a") == "a"`? (false; `.equals` true; `.intern()` fixes)
- `String` vs `StringBuilder` vs `StringBuffer`? When each?
- How many objects does `String s = new String("a")` create? (up to 2: literal in pool + heap object)

7. Interview Traps

- Using `==` to compare string content.
- Saying the pool is in PermGen (it moved to heap in Java 7).
- Using `String +=` in loops.

8. Best Answer

“`String` is immutable and `final` for security, thread-safety, and so it can be pooled and safely cached as a map key. Literals are interned so equal literals share one object, which is why `==` can be true for literals but false against `new String`. For heavy concatenation I use `StringBuilder`; `StringBuffer` only when a buffer is genuinely shared across threads.”

9. Coding Example

```
String a = "hi", b = "hi";
System.out.println(a == b);           // true (pool)
System.out.println(a == new String("hi")); // false (new heap object)
System.out.println(a.equals(new String("hi"))); // true

StringBuilder sb = new StringBuilder();
for (int i = 0; i < 1000; i++) sb.append(i).append(','); // O(n), no garbage storm
String csv = sb.toString();
```

10. Follow-up Coding Questions

- Reverse a string in place using `StringBuilder.reverse()`; then without it.
- Explain why using a mutable object as a `HashMap` key is dangerous.

11 & 12. Summary + Cheat Sheet

Immutable `String` → safe & poolable. `==` = reference, `.equals()` = value. Build with `StringBuilder`; synchronize with `StringBuffer` only if shared.

1.9 Immutable Objects · Wrapper Classes · Autoboxing

Building an immutable class (common coding ask)

```
public final class Money {
    // 1. final class (no subclass)
    private final long cents;
    private final String currency;
    // 2. private final fields
    public Money(long cents, String currency) { // 3. set via constructor
        this.cents = cents;
        this.currency = currency; // 4. defensively copy mutable inputs (none here)
    }
    public long cents() { return cents; } // 5. getters only, no setters
    public String currency() { return currency; }
}
```

Rules: `final` class, all fields `private final`, no setters, defensive copies of mutable fields (arrays, `Date`, collections) on the way in and out.

Wrapper Classes & Autoboxing

- Wrappers box primitives into objects: `int`→`Integer`, `long`→`Long`, etc. Needed for generics/collections (`List<Integer>`).
- **Autoboxing** = automatic `int` → `Integer`; **unboxing** = reverse. Compiler inserts `Integer.valueOf()`/`intValue`.
- **Integer cache**: `valueOf` caches -128..127, so `Integer a=100, b=100; a==b` is `true`, but `127 ... Integer a=200, b=200; a==b` is `false` (different objects). Always compare wrappers with `.equals()`.
- **NPE trap**: unboxing a `null Integer` into an `int` throws `NullPointerException`.

Traps: `Integer ==` comparisons; NPE on unboxing null; autoboxing in tight loops (creates garbage — prefer primitives for performance-critical code).

1.10 Enum · var · Records · Sealed Classes

Enum

Type-safe set of constants; each is a singleton. Can have fields, constructors, methods, and per-constant bodies. Great for strategy/state. Use in `switch`; use `EnumMap`/`EnumSet` (very fast). `enum` is implicitly `final` and extends `java.lang.Enum`.

```
enum Status {
    ACTIVE(1), INACTIVE(0);
    private final int code;
    Status(int code){ this.code = code; }
    int code(){ return code; }
}
```

var (Java 10) — local variable type inference

Only for **local** variables with an initializer. Not for fields, params, or return types. Still statically typed — `var list = new ArrayList<String>();` infers `ArrayList<String>`. Don't overuse; keep readability.

Records (Java 16) — basic awareness

Immutable data carriers: auto-generate constructor, `equals`, `hashCode`, `toString`, accessors.

```
public record Point(int x, int y) {} // that's a full immutable value class
```

Great for DTOs. Can't extend classes; are implicitly `final`. Interviewers love "records vs Lombok `@Value`".

Sealed Classes (Java 17) — basic awareness

Restrict which classes may extend/implement: `sealed interface Shape permits Circle, Square {}`. Enables exhaustive `switch` and controlled hierarchies. Just know the concept and syntax.

Module 1 — Top 25 Interview Questions (with senior answers)

- JDK vs JRE vs JVM?** JDK builds (has `javac`), JRE runs (JVM + libs), JVM executes byte-code. Ship a JRE in prod.
- Is Java compiled or interpreted?** Both — compiled to byte-code, then interpreted + JIT-compiled to native.
- What is JIT / tiered compilation?** Promotes hot methods to native via C1 then C2.
- Class loading phases?** Loading → Linking (verify, prepare, resolve) → Initialization.
- Parent delegation?** Child asks parent first — prevents core-class spoofing & duplicates.
- `ClassNotFoundException` vs `NoClassDefFoundError`?** Reflection-time missing vs runtime-missing after compile.
- Stack vs Heap?** Per-thread frames/locals vs shared objects (GC-managed).
- What replaced PermGen?** Metaspace (native memory, auto-grows) in Java 8.
- How does GC decide liveness?** Reachability from GC roots, not reference counting.
- Generational GC?** Most objects die young → cheap minor GC of Eden; survivors promoted to old gen.
- Default collector? Low-latency choice?** G1 default; ZGC/Shenandoah for sub-ms pauses.
- Can `System.gc()` force GC?** No — only a hint.
- Do cycles leak in Java?** No — reachability handles cycles. Leaks come from unbounded reachable structures.
- `final` vs `finally` vs `finalize`?** Modifier vs cleanup block vs deprecated GC hook.
- Why is `finalize` deprecated?** Unpredictable, no guarantee, hurts GC — use try-with-resources/Cleaner.
- 4 OOP pillars?** Encapsulation, Abstraction, Inheritance, Polymorphism.
- Explain SOLID + a fix.** e.g., Dependency Inversion → inject an interface, not a concrete client.
- Overloading vs overriding?** Compile-time (params) vs runtime (vtable dispatch).
- Can you override `static/private/final`?** No to all — static is hidden; private/final aren't inherited/overridable.
- Interface vs abstract class?** Multiple inheritance + contract vs single inheritance + shared state.
- Why is String immutable?** Security, thread-safety, hashCode caching, pooling.
- `==` vs `.equals()`?** Reference identity vs logical equality.
- String vs StringBuilder vs StringBuffer?** Immutable vs mutable-unsynchronized vs mutable-synchronized.
- Integer caching / autoboxing traps?** -128..127 cached; compare wrappers with `.equals()`; null unboxing NPE.
- What are records/sealed classes for?** Immutable data carriers; restricted hierarchies for exhaustive matching.

Module 1 — Top Coding Questions

- Build an immutable `Money/Employee` class (with defensive copies).
- Implement `equals()` and `hashCode()` correctly (and explain the contract).
- Reverse a string; check palindrome; count char frequency with a `Map`.
- Demonstrate the `Integer` cache boundary (`127` vs `128`).
- Write a class hierarchy showing runtime polymorphism, then refactor an `if/else` type-switch into Strategy.

Module 1 — Common Follow-ups

- “Show me the `equals/hashCode` contract.” (equal objects → equal hash; used by `HashMap`.)
- “How would you debug an OOM in prod?” (heap dump + MAT, GC logs, check caches/`ThreadLocals`.)
- “Why not just use Lombok instead of records?” (records are language-level, immutable, no build-time processor.)

Module 1 — One-Page Cheat Sheet

```
JDK⊃JRE⊃JVM | javac->bytecode->JIT(C1/C2 tiered)
ClassLoader: Load->Verify/Prepare/Resolve->Init | parent-delegation
Heap=objects(GC) Stack=frames(per-thread) Metaspace=class meta(native)
GC: reachability from roots; generational (Eden->Survivor->Old); G1 default, ZGC low-latency
System.gc()=hint. Cycles don't leak. finalize=deprecated -> try-with-resources
OOP: Encapsulation/Abstraction/Inheritance/Polymorphism
SOLID: S O L I D; DI = Dependency Inversion
Overload=compile-time(params); Override=runtime(vtable). static/private/final NOT overridable
Interface: multi-inherit + default methods; Abstract: single + state
String immutable+pooled; ==ref vs .equals value; StringBuilder(fast) StringBuffer(sync)
Integer cache -128..127; compare wrappers with equals; null unbox=NPE
record=immutable DTO; sealed=restricted hierarchy; var=local inference
```

Module 1 — Mock Interview (answer these, then continue)

1. “Walk me through what happens from `javac Foo.java` to the method running as native code.”
2. “My Spring service throws `OutOfMemoryError: Java heap space` at 3 a.m. under load. Walk me through diagnosis and fixes.”
3. “Why can two web apps in one Tomcat use different Jackson versions?”
4. “Give me a real SOLID violation you fixed and how DI helped.”
5. “`Integer a = 128, b = 128; a == b` — true or false, and why? What about `127`?”
6. “Write an immutable `Money` class and justify every keyword.”

Model answers are embedded in the sections above. When ready, continue to Module 2.


Module 2 — Collections Framework

Highest priority and the single most-asked coding-adjacent topic. “How does HashMap work internally?” is asked in *almost every* Java interview at TCS/Infosys/Cognizant and product companies alike. Know the internals cold.

Node.js bridge: JS has `Array`, `Map`, `Set`, `Object`. Java gives you a rich typed hierarchy with explicit complexity guarantees and thread-safe variants. `HashMap` ≈ JS `Map` but with load factor, buckets, and treeification you can be quizzed on.

2.1 Collection Hierarchy

Core Concept



- **Map is not a Collection** — it's a separate root (key → value). Common trap.
- **List** = ordered, indexed, duplicates allowed.
- **Set** = no duplicates.
- **Queue/Deque** = FIFO / double-ended.

Interfaces vs implementations (senior framing)

Program to the interface: `List<X> l = new ArrayList<>();`, `Map<K,V> m = new HashMap<>();`. Choose the implementation for the access pattern and thread model.

Cheat Sheet — pick the right one

Need	Use
Indexed, mostly reads	<code>ArrayList</code>
Frequent add/remove at ends, queue/deque	<code>ArrayDeque</code> (not <code>LinkedList</code>)
Unique, fast lookup	<code>HashSet</code>
Unique + insertion order	<code>LinkedHashSet</code>
Unique + sorted	<code>TreeSet</code>
Key → value fast	<code>HashMap</code>
Key → value + order	<code>LinkedHashMap</code> (also LRU)
Key → value + sorted	<code>TreeMap</code>
Concurrent map	<code>ConcurrentHashMap</code>
Priority ordering	<code>PriorityQueue</code>

2.2 ArrayList — Dynamic Array, Resize, Complexity

1. Why Interviewers Ask This

It's the default `List` and the base for “what's the complexity of `add`” and “how does it grow” questions.

2. Core Concept

Resizable array backed by `Object[] elementData`. Random access by index is $O(1)$.

3. Internal Working — growth

- Default capacity **10** (lazily allocated on first add).
- When full, it grows by **~50%**: `newCap = oldCap + (oldCap >> 1)`. (`ArrayList`; `Vector` doubles.)
- Growth = allocate new array + `System.arraycopy` old → new. That single add is $O(n)$, but **amortized $O(1)$** across many adds.
- `add(index, e)` and `remove(index)` shift elements → $O(n)$.
- Not thread-safe.

4. Memory Diagram

```
size=3, capacity=4: elementData -> [ A | B | C | _ ]
add(D): full -> newcap = 4 + 2 = 6, copy:
           [ A | B | C | D | _ | _ ]
```

5. Real Production Example

Loading 1M rows from DB into `new ArrayList<>()` triggers ~20 resizes/copies. Pre-size with `new ArrayList<>(1_000_000)` to avoid the copy churn — a real latency/GC win.

6. Most Asked Questions

- How does ArrayList grow? (1.5x, `arraycopy`)
- Complexity of get/add/add-at-index/remove? ($O(1)$ /amortized $O(1)$ / $O(n)$ / $O(n)$)
- ArrayList vs Array? (*dynamic vs fixed, generics, autobox overhead*)
- ArrayList vs LinkedList? (see below)
- Is ArrayList thread-safe? How to make it? (`Collections.synchronizedList` / `CopyOnWriteArrayList`)

7. Traps

- Saying it doubles (that's `Vector`; ArrayList is 1.5x).
- Removing while iterating with a for-each → `ConcurrentModificationException` (use `Iterator.remove()`).

8. Best Answer

“`ArrayList` is a dynamic array; get is $O(1)$. When capacity is exceeded it grows by 50% and arraycopies, so add is amortized $O(1)$ but inserts/removes in the middle are $O(n)$ due to shifting. I pre-size it when I know the count to avoid resize churn.”

9. Coding Example

```
List<String> list = new ArrayList<>(1000);    // pre-sized
list.add("a");
// Safe removal during iteration:
Iterator<String> it = list.iterator();
while (it.hasNext()) { if (it.next().isEmpty()) it.remove(); }
```

10. Follow-ups

- Remove all even numbers safely (Iterator vs `removeIf`).

- Convert array ↔ list (`Arrays.asList` gotcha: fixed-size, backed by array).

11 & 12. Summary + Cheat

Dynamic array, 1.5x growth, $O(1)$ get, $O(n)$ middle ops. Pre-size when possible.

2.3 LinkedList

Doubly-linked list; implements `List` and `Deque`. - `get(index) = O(n)` (walks nodes); add/remove at ends = $O(1)$; add/remove in middle = $O(1)$ if you hold the node but $O(n)$ to find it. - Higher memory per element (two pointers + object header). - In practice, prefer `ArrayList` for lists and `ArrayDeque` for queue/stack — `LinkedList` is rarely the right answer despite the textbook “fast inserts”.

Best answer trap: “Use `LinkedList` for frequent insertions” — only true when you already have the node reference; otherwise finding the position is $O(n)$. CPU cache locality makes `ArrayList` win in most real workloads.

2.4 HashMap — the flagship internals question

1. Why Interviewers Ask This

The most-asked Java question, period. It tests hashing, buckets, collisions, equals/hashCode, resize, and (post-Java 8) treeification.

2. Core Concept

Stores key → value in an **array of buckets**. The key's `hashCode()` decides the bucket; `equals()` resolves within a bucket. Average $O(1)$ get/put.

3. Internal Working (Java 8+)

1. `hash(key) = h = key.hashCode(); h ^ (h >>> 16)` — spreads high bits into low bits so poorly-distributed hashcodes still scatter.
2. **Bucket index** = `(n - 1) & hash` where `n` = table length (always a power of 2, so `&` = fast modulo).
3. Each bucket is a **Node linked list**; on collision, append (Java 8 appends at tail).
4. **Put**: if key's hash+`equals` matches an existing node → replace value; else add node.
5. **Load factor** default **0.75**. When `size > capacity * loadFactor`, **resize**: capacity doubles ($16 \rightarrow 32 \rightarrow \dots$), and every node is **rehashed** into the new table.
6. **Treeification**: if a single bucket's list length ≥ 8 and table capacity ≥ 64 , the list converts to a **red-black tree** → worst case $O(\log n)$ instead of $O(n)$. If it shrinks below **6**, it **untreeifies** back to a list.
7. Default initial capacity **16**. Null key allowed (stored in bucket 0); null values allowed.

4. Memory Diagram

```
capacity=16, loadFactor=0.75 -> resize at size 12
table:
[0] -> null
[1] -> (k1,v1) -> (k9,v9)      <- collision: linked list in bucket 1
[2] -> null
...
[5] -> (k5,v5) -> ... (>=8 & cap>=64) -> converts to Red-Black TREE

index = (n-1) & hash(key)      hash(key)=h ^ (h>>>16)
```

5. Real Production Example

Using a mutable object as a key (its `hashCode` changes after insertion) makes entries “disappear” — you can’t `get` them because the bucket index changed. Always use **immutable keys** (`String`, `Long`, records). Overriding `equals` but not `hashCode` breaks map lookups — a classic prod bug.

6. Most Asked Interview Questions

- Explain HashMap internal working. (*hash → bucket → list/tree → equals*)
- What is load factor? Why 0.75? (*space/time tradeoff; higher = more collisions, lower = wasted space*)
- How/when does resize happen? (*size > cap0.75 → double + rehash*)*
- What is treeification and its thresholds? (*list → tree at 8 & cap ≥ 64*)
- Why must capacity be a power of 2? (*(n-1) & hash works as fast modulo*)
- equals/hashCode contract? (*equal → same hash; unequal may collide*)
- What happens with a null key? (*bucket 0*)
- Is HashMap thread-safe? What breaks under concurrency? (*no; resize can corrupt/infinite-loop in Java 7; use CHM*)
- **Follow-up:** difference between Java 7 and 8 HashMap? (*Java 8 added treeification, tail-insert to avoid the Java 7 resize infinite-loop*)

7. Interview Traps

- Overriding `equals` without `hashCode` (or vice versa).
- Saying collisions become a tree immediately (needs length ≥ 8 **and** capacity ≥ 64).
- Saying HashMap maintains order (it doesn’t — use `LinkedHashMap`).
- Claiming HashMap put/get is always O(1) (worst case O(log n) after treeify, O(n) before Java 8).

8. Best Answer

“HashMap is an array of buckets. I compute a spread hash (`h ^ h >>> 16`) and index with `(n-1) & hash`. Collisions chain in a linked list; from Java 8, once a bucket has ≥ 8 nodes and the table is ≥ 64, it becomes a red-black tree for O(log n) worst case. It resizes — doubling capacity and rehashing — when size exceeds capacity × 0.75. Keys must be immutable with a consistent `equals` / `hashCode`, or lookups break.”

9. Coding Example — correct equals/hashCode

```
public final class UserId {
    private final String tenant;
    private final long id;
    public UserId(String tenant, long id){ this.tenant = tenant; this.id = id; }

    @Override public boolean equals(Object o) {
        if (this == o) return true;
        if (!(o instanceof UserId u)) return false;
        return id == u.id && tenant.equals(u.tenant);
    }
    @Override public int hashCode() { return Objects.hash(tenant, id); }
}
// Now safe as a HashMap key.
```

10. Follow-up Coding Questions

- Find the first non-repeating char using a `LinkedHashMap<Character, Integer>`.
- Group anagrams using `HashMap<String, List<String>>`.
- Implement a simple hash map from scratch (array + linked list + resize).

11. Summary

Array of buckets, spread hash, `(n-1) & hash` index, chaining → tree at 8/64, resize at 0.75. Immutable keys + correct equals/hashCode.

12. Cheat Sheet

```
default cap 16, LF 0.75, resize=double+rehash
index=(n-1)&hash ; hash=h^(h>>16)
treeify: bucket>=8 AND cap>=64 ; untreeify<6
null key -> bucket 0 ; not thread-safe
equals true => hashCode equal (mandatory)
```

2.5 ConcurrentHashMap (CHM)

Core Concept & Internal Working

Thread-safe **Map** without locking the whole table. - **Java 7**: segment locking (default 16 segments) — lock striping. - **Java 8+**: dropped segments. Uses the **bucket array + CAS** for empty-bucket inserts and **synchronized on the bucket's head node** for collisions. Reads are lock-free (**volatile** nodes). Also treeifies like **HashMap**. - **No null keys or values** (ambiguity between “absent” and “null” in concurrent **get**). - **size()** is approximate under concurrency (uses **baseCount** + counter cells). - Atomic composites: **putIfAbsent**, **compute**, **computeIfAbsent**, **merge** — use these instead of check-then-act.

Memory Diagram

```
Java 8 CHM:
bucket empty -> CAS in new node (no lock)
bucket has head-> synchronized(head){ append / update } (fine-grained)
reads -> volatile, lock-free
```

Interview Q&A

- **HashMap vs Hashtable vs CHM?** **HashMap** not thread-safe; **Hashtable** synchronizes every method (whole-object lock, legacy, slow); **CHM** locks per-bucket → high concurrency.
- **Why no null in CHM?** Can't distinguish “key absent” from “mapped to null” without locking.
- **Why not just Collections.synchronizedMap?** That wraps a single lock — no concurrency; **CHM** allows concurrent reads and striped writes.
- **computeIfAbsent use?** Atomic lazy init of cache entries (avoids race of **if(!contains) put**).

Best Answer

“CHM gives thread-safety with high throughput: in Java 8 it CASes into empty buckets and only **synchronized** - locks the individual bucket head on collision, while reads stay lock-free via **volatile**. It forbids nulls and offers atomic **computeIfAbsent/merge** so I don't do racy check-then-act. **Hashtable** locks the whole map — I never use it.”

2.6 Sets — HashSet, LinkedHashSet, TreeSet

	HashSet	LinkedHashSet	TreeSet
Backing	HashMap	LinkedHashMap	TreeMap (red-black)
Order	none	insertion	sorted
Null	1 null	1 null	no null (NPE on compare)
get/add/contains	O(1)	O(1)	O(log n)

- **HashSet** is literally a **HashMap** with a dummy **PRESENT** value for every key. Same equals/hashCode rules apply.

- **TreeSet** needs elements **Comparable** or a **Comparator**; keeps sorted order; offers **first/last/ceiling/floor/headSet/tailSet**.
- **LinkedHashSet** = predictable iteration order (insertion), slightly more memory.

2.7 Queues & Deque — PriorityQueue, ArrayDeque

PriorityQueue


- A **binary min-heap** (array-backed). **peek** = min = $O(1)$; **offer/poll** = $O(\log n)$.
- Ordering by natural order or a **Comparator**. **Not** sorted iteration — only the head is guaranteed smallest.
- Not thread-safe (use **PriorityBlockingQueue**).
- Classic uses: Dijkstra, top-K, task scheduling by priority.

ArrayDeque

- Resizable-array double-ended queue. **Preferred stack and queue** (faster than **Stack** and **LinkedList**; no synchronization overhead).
- **offerFirst/offerLast/pollFirst/pollLast**, or **push/pop** for stack semantics. No null allowed.

Memory Diagram — PriorityQueue heap

Array: [1, 3, 2, 7, 4] represents min-heap:



parent(i)=(i-1)/2 ; children=2i+1, 2i+2 ; poll() removes root, sifts down

2.8 Maps — TreeMap, LinkedHashMap

TreeMap

- **Red-black tree**, sorted by key (natural or **Comparator**). $O(\log n)$ ops.
- **NavigableMap**: **firstKey/lastKey/ceilingKey/floorKey/headMap/tailMap/subMap**.
- No null keys. Use when you need **range queries** or sorted iteration.

LinkedHashMap

- HashMap + doubly-linked list across entries → predictable **insertion** (or **access**) order.
- **accessOrder=true** + **override removeEldestEntry** → a ready-made **LRU cache** (top interview coding ask).

```

class LRUCache<K,V> extends LinkedHashMap<K,V> {
    private final int cap;
    LRUCache(int cap){ super(cap, 0.75f, true); this.cap = cap; } // access-order
    @Override protected boolean removeEldestEntry(Map.Entry<K,V> e){ return size() > cap; }
}
  
```

2.9 Comparable vs Comparator

	Comparable	Comparator
Method	<code>compareTo(o)</code>	<code>compare(a,b)</code>
Where	inside the class	external/separate
Count	one “natural” order	many orders
Example	<code>String</code> , <code>Integer</code>	sort users by age then name

```
class Emp { String name; int age; }
List<Emp> list = ...;
list.sort(Comparator.comparingInt((Emp e) -> e.age) // Comparator, Java 8 style
    .thenComparing(e -> e.name));
// reverse: .reversed()
```

- `compareTo/compare` return negative/zero/positive.
- **Trap:** `compareTo` inconsistent with `equals` breaks `TreeSet/TreeMap` (they use compare, not equals — two “equal by compare” elements are deduped!).

2.10 Iterator · Fail-Fast vs Fail-Safe · ConcurrentModificationException

Core Concept

- **Iterator** = cursor to traverse a collection; supports safe `remove()`.
- **Fail-fast** iterators (`ArrayList`, `HashMap`) track a `modCount`. If the collection is structurally modified during iteration (except via the iterator), the next `next()` throws `ConcurrentModificationException` (CME) — best-effort, not guaranteed.
- **Fail-safe** iterators (`CopyOnWriteArrayList`, `ConcurrentHashMap`) iterate over a **snapshot** or tolerate concurrent modification → no CME, but may not reflect latest writes.

Internal Working — modCount

```
expectedModCount = modCount (at iterator creation)
list.add(x) during loop -> modCount++
it.next() -> if (modCount != expectedModCount) throw ConcurrentModificationException
```

CME can happen in a **single thread** too — e.g., adding/removing inside a for-each loop.

Best Answer & Fixes

“Fail-fast iterators throw `ConcurrentModificationException` when they detect structural change via `modCount`, even single-threaded — like removing inside a for-each. I fix it with `Iterator.remove()`, `removeIf()`, collecting to a new list, or a concurrent/copy-on-write collection when I truly have multiple threads.”

```
// WRONG - throws CME:
for (String s : list) if (s.isBlank()) list.remove(s);
// RIGHT:
list.removeIf(String::isBlank);
```

2.11 Collections & Arrays Utility Classes

Collections (operates on Collection objects)

`sort`, `reverse`, `shuffle`, `binarySearch`, `max/min`, `frequency`, `unmodifiableList`, `synchronizedList/Map`, `emptyList`, `singletonList`.

Arrays (operates on arrays)

`sort`, `binarySearch`, `fill`, `copyOf`, `equals`, `asList`, `stream`, `toString`, `parallelSort`. - **Trap:** `Arrays.asList(arr)` returns a **fixed-size** list backed by the array — `add/remove` throw `UnsupportedOperationException`; changes write through to the array. For a mutable copy: `new ArrayList<>(Arrays.asList(arr))`. - `List.of(...)` / `Map.of(...)` (Java 9) → **immutable**; `add` throws.

2.12 Time & Space Complexity (must-memorize table)

Structure	get/contains	add	remove	Notes
ArrayList	O(1) index / O(n) search	amortized O(1) end, O(n) middle	O(n)	array copy on resize
LinkedList	O(n)	O(1) ends	O(1) w/ node	2 pointers/node
ArrayDeque	O(n) search	O(1) ends	O(1) ends	best stack/queue
HashMap/HashSet	O(1) avg, O(log n) treeified	O(1) avg	O(1) avg	0.75 LF
LinkedHashMap/Set	O(1)	O(1)	O(1)	keeps order
TreeMap/TreeSet	O(log n)	O(log n)	O(log n)	sorted, red-black
PriorityQueue	O(1) peek, O(n) contains	O(log n)	O(log n)	binary heap
ConcurrentHashMap	O(1) avg	O(1) avg	O(1) avg	lock-stripped/CAS

Module 2 — Top 25 Interview Questions (senior answers)

1. **Is Map a Collection?** No — separate hierarchy.
2. **How does HashMap work internally?** hash → $(n-1) \& \text{hash}$ bucket → list/tree → equals.
3. **Load factor & default capacity?** 0.75, 16; resize doubles at $\text{size} > \text{cap} * 0.75$.
4. **Treeification thresholds?** ≥ 8 in a bucket and $\text{cap} \geq 64$ → red-black tree; < 6 untreeify.
5. **Why capacity power of 2?** $(n-1) \& \text{hash}$ acts as fast modulo and distributes evenly.
6. **equals/hashCode contract?** Equal objects must have equal hashCode; needed for map/set correctness.
7. **What if you override equals but not hashCode?** Lookups fail — objects land in wrong/duplicate buckets.
8. **HashMap vs Hashtable vs ConcurrentHashMap?** Unsynchronized vs whole-object-locked vs bucket-level/CAS.
9. **Why no null in ConcurrentHashMap?** Can't distinguish absent vs null concurrently.
10. **ArrayList vs LinkedList?** Array O(1) get, cache-friendly) vs nodes O(1) ends only).
11. **How does ArrayList grow?** $1.5x + \text{arraycopy}$; amortized O(1) add.
12. **Vector vs ArrayList?** Vector synchronized + doubles; ArrayList not synced + $1.5x$.
13. **HashSet vs TreeSet vs LinkedHashSet?** Unordered O(1) / sorted O(log n) / insertion-order.
14. **Comparable vs Comparator?** Natural single order in-class vs multiple external orders.
15. **compareTo inconsistent with equals — impact?** TreeSet/TreeMap drop "equal" elements.
16. **Fail-fast vs fail-safe?** modCount CME vs snapshot iteration.
17. **When does ConcurrentModificationException happen?** Structural change during iteration (even single-thread).
18. **How to remove during iteration?** `Iterator.remove()` / `removeIf`.

19. **PriorityQueue internals?** Binary min-heap; $O(\log n)$ offer/poll.
20. **Best stack/queue implementation?** `ArrayDeque` (not `Stack/LinkedList`).
21. **How to build an LRU cache?** `LinkedHashMap` access-order + `removeEldestEntry`.
22. **`Arrays.asList` gotcha?** Fixed-size, array-backed; wrap in `new ArrayList<>()`.
23. **Immutable collections?** `List.of`, `Map.of`, `Collections.unmodifiableList`.
24. **When `TreeMap` over `HashMap`?** Sorted keys / range queries (`ceilingKey`, `subMap`).
25. **Complexity of `HashMap` get worst case?** $O(\log n)$ after treeify ($O(n)$ pre-Java 8).

Module 2 — Top Coding Questions

- Implement LRU cache (`LinkedHashMap` and manual `HashMap`+DLL).
- First non-repeating character (`LinkedHashMap`).
- Group anagrams; find duplicates; count word frequency.
- Top-K frequent elements (`HashMap` + `PriorityQueue`).
- Merge/sort by multiple fields with `Comparator.comparing().thenComparing()`.
- Detect and remove duplicates preserving order (`LinkedHashSet`).
- Implement your own `HashMap` (buckets + resize).

Module 2 — Common Follow-ups

- “Java 7 vs Java 8 `HashMap`?” (treeify + tail-insert fixed resize infinite loop.)
- “Why did old `HashMap` infinite-loop under concurrency?” (Java 7 head-insert reversed list on concurrent resize.)
- “How is CHM `size()` computed?” (`baseCount` + `CounterCells`; approximate.)
- “What happens to iteration order after resize?” (rehash changes bucket order; never rely on `HashMap` order.)

Module 2 — One-Page Cheat Sheet

```
Map != Collection. List(dup,ordered) Set(unique) Queue(FIFO) Deque(both)
ArrayList: dynamic array, 1.5x grow, O(1) get, O(n) middle
LinkedList: doubly-linked, O(1) ends; prefer ArrayList/ArrayDeque
HashMap: cap16 LF0.75, index=(n-1)&hash, hash=h^h>>>16, chain->tree(8 & cap>=64), resize=double+rehash
    immutable keys + equals/hashCode required; null key->bucket0
CHM: Java8 CAS empty + synchronized head; no nulls; computeIfAbsent atomic
Sets: HashSet=HashMap; LinkedHashSet=order; TreeSet=sorted O(log n)
Maps: TreeMap sorted(rbtrees); LinkedHashMap=order/LRU(accessOrder+removeEldestEntry)
PriorityQueue=min-heap O(log n); ArrayDeque=best stack/queue
Comparable=natural(compareTo); Comparator=custom(compare, thenComparing, reversed)
fail-fast(modCount->CME) vs fail-safe(snapshot). remove via Iterator.remove/removeIf
Arrays.asList = fixed-size backed array; List.of = immutable
```

Module 2 — Mock Interview (answer, then continue)

1. “Walk me through `map.put(key, value)` end to end, including collision and resize.”
2. “You override `equals()` on an entity but forget `hashCode()`. What breaks and why?”
3. “How would you build an LRU cache in Java? Now make it thread-safe.”
4. “Explain how `ConcurrentHashMap` achieves thread-safety without locking the whole map.”
5. “Why is `ArrayDeque` preferred over `Stack` and `LinkedList`?”
6. “You get a `ConcurrentModificationException` while filtering a list in one thread. Why, and 3 ways to fix it?”

Model answers are in the sections above. Continue to Module 3 when ready.

Module 3 — Java 8+ (Lambdas, Functional Interfaces, Streams, Optional)

Every modern Java interview has a “write this with streams” coding round. Service companies ask stream/lambda syntax; product companies ask stream *internals* and lazy evaluation.

Node.js bridge: Java streams ≈ JS array methods (`map`, `filter`, `reduce`) but **lazy** and one-shot. Lambdas ≈ arrow functions. Functional interfaces ≈ callback type signatures.

3.1 Lambdas & Functional Interfaces

1. Why Interviewers Ask This

Foundation for streams and modern Spring (functional endpoints, `Optional`, callbacks).

2. Core Concept

- A **functional interface** = exactly one abstract method (`@FunctionalInterface`). May have `default/static` methods.
- A **lambda** is a concise implementation of that single method: `(args) -> body`.
- Method reference** = shorthand for a lambda that just calls an existing method: `String::toUpperCase`.

3. Internal Working

Lambdas are **not** anonymous inner classes. `javac` emits an `invokedynamic` bytecode; at first execution `LambdaMetafactory` spins up the implementation class (often reusing a singleton for stateless lambdas). Result: lower overhead and no extra `.class` file per lambda. A lambda captures **effectively final** local variables (copied), and `this` refers to the **enclosing** instance (unlike anonymous classes).

4. Memory Diagram

```
Runnable r = () -> log(x); // x must be effectively final (captured by value)
javac -> invokedynamic -> LambdaMetafactory -> synthetic impl (cached if stateless)
'this' inside lambda == enclosing object (NOT a new anonymous instance)
```

5. The 4 core functional interfaces (memorize)

Interface	Signature	Meaning	Example
<code>Function<T,R></code>	<code>R apply(T)</code>	transform	<code>s -> s.length()</code>
<code>Predicate<T></code>	<code>boolean test(T)</code>	condition	<code>s -> s.isEmpty()</code>
<code>Consumer<T></code>	<code>void accept(T)</code>	side effect	<code>System.out::println</code>
<code>Supplier<T></code>	<code>T get()</code>	produce	<code>() -> new ArrayList<>()</code>

Also: `BiFunction`, `BinaryOperator`, `UnaryOperator`, `BiPredicate`, `BiConsumer`, and primitive variants (`IntFunction`, `ToIntFunction`, `IntPredicate`) to avoid boxing.

6. Most Asked Questions

- What is a functional interface? Name the built-in ones.
- Lambda vs anonymous class? (*this scope, invokedynamic, no separate class*)
- What is “effectively final”? (*captured locals can't be reassigned*)
- 4 kinds of method references? (*static `C::m`, instance-of-object `obj::m`, instance-of-type `C::m`, constructor `C::new`*)

7. Traps

- Thinking a lambda's `this` = the lambda (it's the enclosing object).
- Trying to mutate a captured local (compile error).
- Overusing lambdas where a named method is clearer.

8. Best Answer

"A functional interface has one abstract method; a lambda implements it and compiles to `invokedynamic` via `LambdaMetafactory`, not an inner class — so `this` is the enclosing object and stateless lambdas are cached. Captured locals must be effectively final. I lean on `Function`, `Predicate`, `Consumer`, `Supplier` and method references for readability."

9. Coding Example

```
Function<String,Integer> len = String::length;           // method reference
Predicate<String> nonEmpty = s -> !s.isEmpty();
Supplier<List<String>> newList = ArrayList::new;         // constructor ref
Consumer<String> print = System.out::println;
BiFunction<Integer,Integer,Integer> add = Integer::sum;
System.out.println(len.andThen(n -> n * 2).apply("hello")); // 10
```

10. Follow-ups

- Compose predicates with `.and()/.or()/.negate()`.
- Write a generic `retry(Supplier<T>, int)` helper.

11 & 12. Summary + Cheat

FI = 1 abstract method; lambda = its impl via `invokedynamic`. `Function`/`Predicate`/`Consumer`/`Supplier`.

3.2 Streams — the flagship Java 8 topic

1. Why Interviewers Ask This

"Rewrite this loop with streams" and "what does this stream print" appear in nearly every coding round.

2. Core Concept

A **Stream** is a pipeline over a data source: **source** → **intermediate ops** → **terminal op**. Streams don't store data, don't mutate the source, and are **single-use** (consumed once).

3. Internal Working — laziness & fusion

- **Intermediate ops** (`map`, `filter`, `sorted`, `distinct`, `limit`, `flatMap`, `peek`) are **lazy** — they build a pipeline and do nothing until a terminal op runs.
- **Terminal ops** (`collect`, `forEach`, `reduce`, `count`, `findFirst`, `anyMatch`, `toList`) trigger execution.
- Elements are pushed **one at a time** through the whole chain (**loop fusion**) — not materialized between stages. This enables **short-circuiting** (`findFirst`, `limit`, `anyMatch` stop early) and lazy `map` calls that never run for skipped elements.
- `parallelStream()` splits work across the common `ForkJoinPool`. Use only for large, CPU-bound, stateless, associative work — and never mutate shared state.

4. Memory Diagram

```
source -> filter -> map -> collect
[1,2,3,4] each element flows fully through the chain before the next:
1 -> filter(even?) drop
2 -> filter ok -> map(*10)=20 -> collector
3 -> drop
4 -> ok -> 40 -> collector      => [20,40]
Nothing runs until collect() (terminal). limit(1) would stop after first match.
```

5. Key operations

- `map` — 1:1 transform. `flatMap` — 1:many, flattens nested streams (`List<List<T>>` → `Stream<T>`).
- `filter` — keep matching. `reduce` — fold to a single value.
- `collect(Collectors...)` — `toList`, `toSet`, `toMap`, `joining`, `groupingBy`, `partitioningBy`, `counting`, `summingInt`, `averagingDouble`.
- `groupingBy` — `Map<K, List<V>>` (SQL GROUP BY). `partitioningBy` — `Map<Boolean, List<V>>` (split by predicate).

6. Most Asked Questions

- Intermediate vs terminal ops? Give examples.
- Are streams lazy? What triggers execution?
- `map` vs `flatMap`?
- Can you reuse a stream? (No — `IllegalStateException`.)
- When parallel streams? Risks? (*shared mutable state, small data overhead, ordering.*)
- `Collectors.toMap` duplicate key behavior? (*throws unless you pass a merge function.*)

7. Traps

- Reusing a consumed stream.
- Side effects in `map/peek` (should be pure).
- `toMap` without a merge function on duplicate keys → `IllegalStateException`.
- Assuming parallel streams are always faster (usually slower for small/I/O work).
- `forEach` on a parallel stream expecting order (use `forEachOrdered`).

8. Best Answer

“A stream is a lazy pipeline: intermediate ops just describe work and only run when a terminal op pulls elements through, one at a time with loop fusion, which lets short-circuit ops like `findFirst` stop early. `map` is 1:1, `flatMap` flattens nested structures. I use `groupingBy/partitioningBy` for aggregation and avoid parallel streams unless the workload is large, CPU-bound, and stateless.”

9. Coding Example

```
record Emp(String name, String dept, int salary) {}
List<Emp> emps = ...;

// group salaries by department
Map<String, List<Emp>> byDept = emps.stream()
    .collect(Collectors.groupingBy(Emp::dept));

// average salary per department
Map<String, Double> avg = emps.stream()
    .collect(Collectors.groupingBy(Emp::dept, Collectors.averagingInt(Emp::salary)));

// highest paid per department
Map<String, Optional<Emp>> top = emps.stream()
    .collect(Collectors.groupingBy(Emp::dept,
        Collectors.maxBy(Comparator.comparingInt(Emp::salary))));

// names of employees earning > 100k, sorted, comma-joined
String csv = emps.stream()
    .filter(e -> e.salary() > 100_000)
    .map(Emp::name).sorted()
    .collect(Collectors.joining(", "));

// partition into high/low earners
Map<Boolean, List<Emp>> parts = emps.stream()
    .collect(Collectors.partitioningBy(e -> e.salary() > 100_000));

// flatMap: all distinct chars across names
List<Character> chars = emps.stream()
    .flatMap(e -> e.name().chars().mapToObj(c -> (char) c))
    .distinct().toList();
```

10. Follow-up Coding Questions

- Frequency count with `Collectors.groupingBy(x, counting())`.
- Second highest salary using streams.
- Sum of squares of even numbers with `reduce/sum`.
- Flatten `List<List<Integer>>` to a sorted distinct list.
- Convert `List<Emp>` to `Map<name, salary>` handling duplicate names.

11 & 12. Summary + Cheat

Lazy pipeline, single-use, terminal triggers. `map/flatMap/filter/reduce/collect`; `groupingBy/partitioningBy` for aggregation.

3.3 Optional

Core Concept & Internal Working

A container that may or may not hold a non-null value — designed to make “absence” explicit at the type level and reduce NPEs. Intended primarily as a **return type**, not for fields or parameters.

API you must know

`Optional.of(x)` (x must be non-null), `ofNullable(x)`, `empty()`, `isPresent()/isEmpty()`, `get()` (avoid), `orElse(default)`, `orElseGet(supplier)`, `orElseThrow()`, `map`, `flatMap`, `filter`, `ifPresent(consumer)`, `ifPresentOrElse(...)`.

Best practices / traps

- **Never** call `get()` without checking — defeats the purpose and NPE-equivalent.
- **orElse vs orElseGet**: `orElse(expensive())` **always** evaluates the argument even when present; `orElseGet() -> expensive()` is lazy. Use `orElseGet` for costly defaults.

- Don't use `Optional` for fields, method params, or collections (use empty collections instead).
- Don't do `if (opt.isPresent()) opt.get()` — use `map/flatMap/orElse`.

```
// Spring Data returns Optional:
User user = repo.findById(id)
    .orElseThrow(() -> new NotFoundException("user " + id));

String city = repo.findById(id)
    .map(User::getAddress)
    .map(Address::getCity)
    .orElse("UNKNOWN");           // null-safe chain, no NPE
```

Best Answer

“`Optional` makes absence explicit in the return type so callers must handle it. I chain `map/flatMap` for null-safe navigation and use `orElseThrow/orElseGet` — never bare `get()`. I avoid it for fields and parameters; that's an anti-pattern.”

3.4 Method References (quick reference)

Kind	Syntax	Lambda equivalent
Static	<code>Integer::parseInt</code>	<code>s -> Integer.parseInt(s)</code>
Instance of a <i>particular</i> object	<code>System.out::println</code>	<code>x -> System.out.println(x)</code>
Instance of an <i>arbitrary</i> object of a type	<code>String::toUpperCase</code>	<code>s -> s.toUpperCase()</code>
Constructor	<code>ArrayList::new</code>	<code>() -> new ArrayList<>()</code>

Module 3 — Top 25 Interview Questions (senior answers)

1. **Functional interface?** One abstract method; `@FunctionalInterface`.
2. **Built-in functional interfaces?** Function, Predicate, Consumer, Supplier (+ Bi/primitive variants).
3. **Lambda vs anonymous class?** invokedynamic, enclosing `this`, no extra class.
4. **Effectively final?** Captured locals can't be reassigned.
5. **Method reference kinds?** static, bound instance, unbound instance, constructor.
6. **What is a Stream?** Lazy single-use pipeline over a source.
7. **Intermediate vs terminal ops?** Lazy builders vs execution triggers.
8. **Is stream lazy? Proof?** Yes — `peek/map` don't run until terminal; short-circuits stop early.
9. **map vs flatMap?** 1:1 vs 1:many flatten.
10. **reduce use?** Fold to single value with identity + accumulator (+ combiner for parallel).
11. **collect / Collectors?** `toList/toMap/joining/groupingBy/partitioningBy/counting`.
12. **groupingBy vs partitioningBy?** Arbitrary key map vs boolean 2-way split.
13. **Can you reuse a stream?** No — `IllegalStateException`.
14. **Parallel stream risks?** Shared mutable state, overhead, ordering, common pool starvation.
15. **toMap duplicate keys?** Throws unless merge function provided.
16. **Stream vs Collection?** Computation pipeline vs data storage.
17. **Optional purpose?** Explicit absence, fewer NPEs; return type only.
18. **orElse vs orElseGet?** Eager arg vs lazy supplier.
19. **Why not Optional fields/params?** Overhead + not designed for it; use empty collections.
20. **findFirst vs findAny?** Deterministic first vs any (parallel-friendly).

21. **peek use/abuse?** Debugging only; not for side effects driving logic.
22. **mapToInt/IntStream?** Avoid boxing; get `sum/average/summaryStatistics`.
23. **Collectors.joining?** Concatenate with delimiter/prefix/suffix.
24. **How does parallelStream split?** Spliterator + common ForkJoinPool.
25. **Infinite streams?** `Stream.iterate/generate` + `limit` (must short-circuit).

Module 3 — Top Coding Questions

- Group employees by dept; average/max salary per dept.
- Find duplicates / first non-repeating char with streams.
- Second highest number; top-K frequent words.
- Flatten nested lists (`flatMap`); merge maps.
- Word frequency count; sort a `Map` by value.
- Partition numbers into even/odd; sum with `reduce`.

Module 3 — Common Follow-ups

- “Show it with a `for` loop, then streams — which is clearer here?”
- “Would you parallelize this? Why/why not?”
- “How does `Collectors.toMap` behave on collisions and how do you fix it?”

Module 3 — One-Page Cheat Sheet

```
Functional interface = 1 abstract method. Lambda -> invokedynamic (not inner class), enclosing this
Function(apply) Predicate(test) Consumer(accept) Supplier(get)
Method refs: Class::static, obj::method, Class::instanceMethod, Class::new
Stream = lazy, single-use pipeline. Intermediate(map/filter/flatMap/sorted/limit) lazy;
Terminal(collect/reduce/forEach/count/findFirst/anyMatch) triggers.
map=1:1, flatMap=flatten. Collectors: toList/toMap(merge!)/joining/groupingBy/partitioningBy/counting
Optional: return type only; map/flatMap/orElseThrow; orElse(eager) vs orElseGet(lazy); never bare get()
parallelStream only for big CPU-bound stateless work
```

Module 3 — Mock Interview (answer, then continue)

1. “Given `List<Order>`, produce total revenue per customer, sorted descending, as a `LinkedHashMap`.”
2. “Explain, step by step, what elements flow through `stream().filter(...).map(...).findFirst()` and why `map` may run fewer times than the list size.”
3. “`orElse` vs `orElseGet` — when does it matter in production?”
4. “Rewrite a nested loop that flattens and dedups a list of lists using streams.”
5. “When would you refuse to use a parallel stream?”

Continue to Module 4 when ready.